

Per-sample metrics vs. inference compute
 GPT-5.4, cyber, bg=100× (filled = main, faint = resample)

